

SAFE STATIC INITIALISATION AND CLEANUP IN LIBRARIES

ASHLEY ROLL

 Code and Slides: github.com/AshleyRoll/conferences/cppnorth2024

TOPICS

- Storage Duration and Linkage
- Non-local and local variable initialisation
- How Static Storage Duration is managed by the compiler
- Common Static initialisation patterns
- Libraries and Static Storage initialisation
- Walk-through example library code

ASSUMPTIONS

- C++17 and above
- Not covering `thread_local` variables / `thread` storage duration
- Not covering runtime load and unload of shared libraries
- Trying to strike a balance between *understanding* and *exact*
 - some implementation details glossed over

WHAT IS STATIC INITIALIZATION?

- Storage Duration for values that live outside a scope
- Often Global State
 - Logger
 - Thread pools
 - Resource managers
- Iostream has `std::cin`, `std::cout`, etc
- 3rd Party libraries like Curl and OpenSSL can have startup and shutdown requirements

STORAGE DURATION

- *automatic* - **stack**
block entered → block exited
- *static* - **fixed memory**
program begins → program ends
- *dynamic* - **heap**
new → **delete**

```
1 static int localCounter{0}; // static duration
2 int globalCounter{10};      // static duration
3
4 namespace { // anonymous namespace
5     int localValue{1};      // static duration
6 }
7
8 namespace ns {
9     int localValue{2};      // static duration
10 }
11
12 void function() {
13     int i{10};              // automatic duration
14     static int j{0};        // static duration
15     int *a =                 // "a" automatic duration
16         new int[10];        // "object" dynamic duration
17     delete [] a;
18 }
19
20 struct jar {
21     static int pickle_count;
22 };
23 // define once in program
24 int jar::pickle_count{10}; // static duration
25
26 struct bottle {
27     // many definitions, linker chooses one
28     inline static double litres{0.5}; // static duration
29 };
```

LINKAGE

- *No Linkage*
only the current scope can access
- *Internal Linkage*
only the current translation unit can access
- *External Linkage*
all translation units can access single instance

A variable with *Linkage* has a fixed location in the program allocated by the linker. This is done using segments (*.bss*, *.data*, *.rodata*)

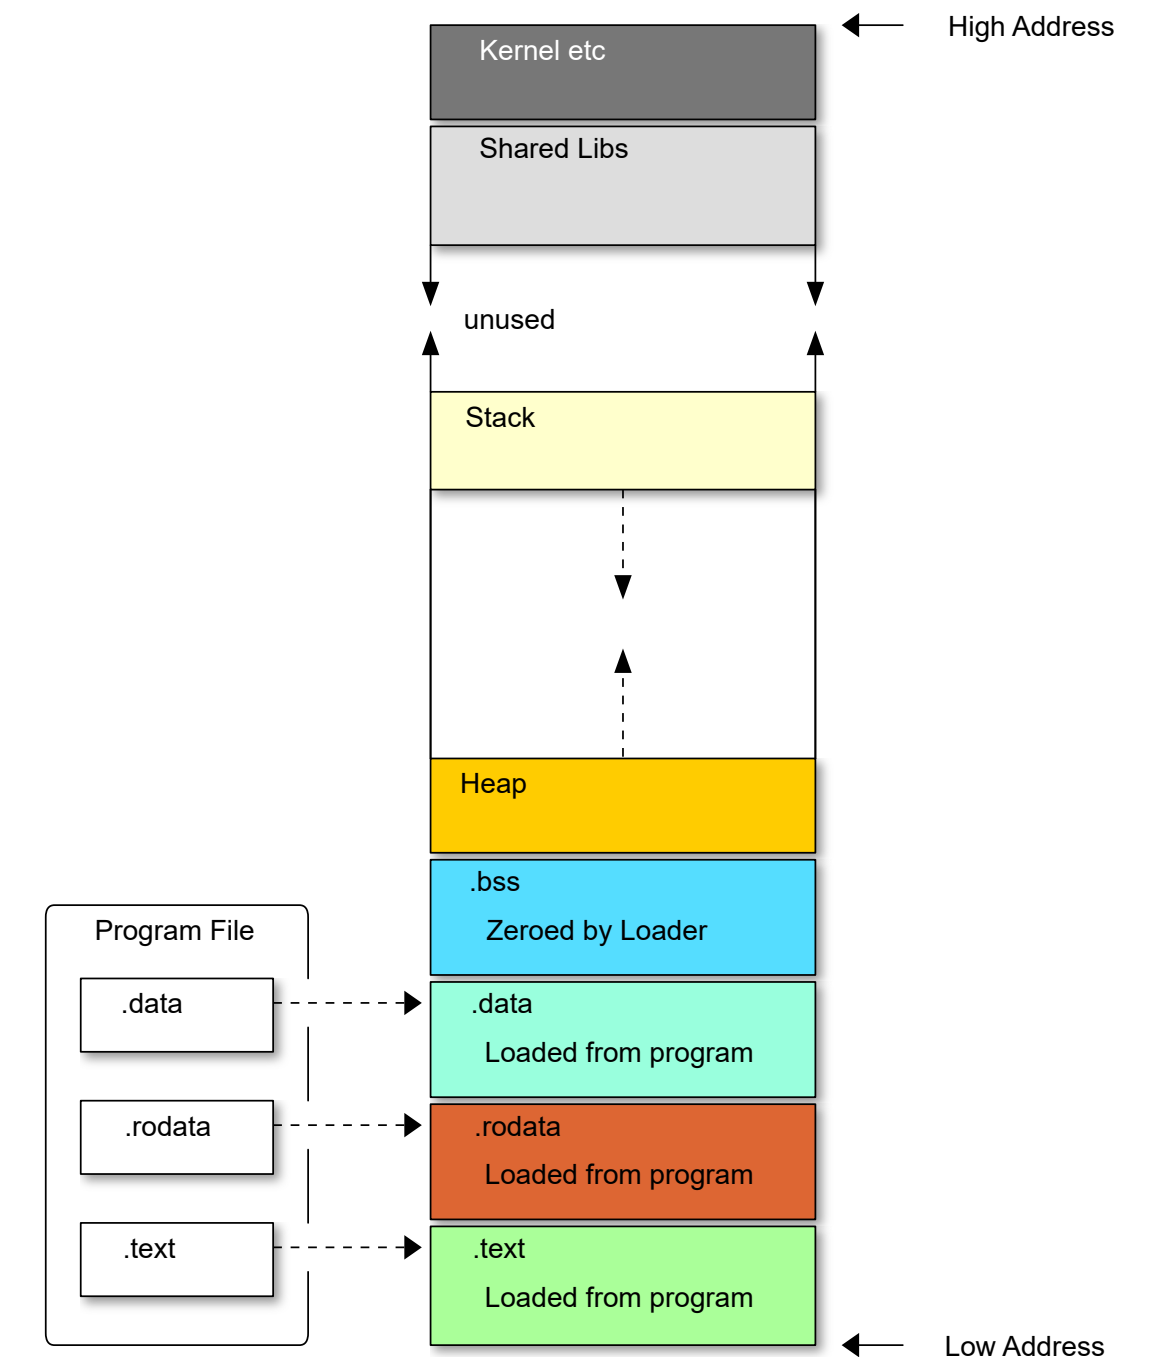
```
1 static int localCounter{0}; // internal linkage
2 int globalCounter{10};      // external linkage
3
4 namespace { // anonymous namespace
5     int value{1};            // internal linkage
6 }
7
8 namespace ns {
9     int globalValue{2};      // external linkage
10    static int value{2};      // internal linkage
11 }
12
13
14 void function() {
15     int i{10};               // no linkage
16     static int j{0};         // no linkage
17     int *a =                  // "a" no linkage
18         new int[10];         // "object" no linkage
19     delete [] a;
20 }
21
22 struct jar {
23     static int pickle_count;
24 };
25 // define once in program
26 int jar::pickle_count{10};    // external linkage
27
28 struct bottle {
29     // many definitions, linker chooses one
30     inline static double litres{0.5}; // external linkage
31 };
```

DETOUR: PROGRAM LOADING AND RUNNING

- Loader & Kernel setup memory space
- Loader places program segments in memory
- Loader links any shared libraries in
- Loader establishes **Heap** and **Stack**
- Runs program (jumps to **_start** in runtime lib)
 - Initialisation routines run
 - **main()** called
 - Finaliser routines run
 - **exit**



Slideware - some simplifications



NON-LOCAL VARIABLES

- Scoped at namespace level, including global namespace
- `c`lass or `s`truct level static variables, including templates
- Have *static* storage duration
- Have *internal* or *external* linkage
- Mapped into fixed location in `.data` (mutable) / `.rodata` (const) or `.bss` sections
- ⚠ *In Principle*^[1] ⚠ initialised before program starts

1. Your Mileage May Vary, what most of the rest of this talk is about

LOCAL VARIABLES OF STATIC STORAGE DURATION

- `static` variables inside block scopes
- Have *static* storage duration
- Have *no* linkage
- Mapped into fixed location in `.data` / `.rodata` or `.bss` sections
- Initialised first time control flow passes their declaration^[1]
 - Compiler must ensure this is thread-safe
 - Has some minor cost associated each time control passes through
 - Compiler can move initialisation to before program start in some cases

1. We will also explore this further

STATIC STORAGE INITIALIZATION IN THE C++ STANDARD

- Static Initialisation
 - Constant initialisation
 - Zero initialisation
- Dynamic Initialisation
- Static local variables

CONSTANT INITIALISATION

- Compile time known constant (not zero)
- Placed into `.data` (or `.rodata`) segments

`basic.start.static`

ZERO INITIALISATION

- IF Compile time known to be zero (`0`, `nullptr`)
- Placed into `.bss` segment

`basic.start.static`

DYNAMIC INITIALISATION

- Everything else
- Placed into `.bss` segment (zeroed); **AND** *Dynamic initialisation* code emitted to compute value
- Ordering is complex:
 - **Unordered** - class template static members
 - **Partially Ordered** - inline static variables IFF consistently ordered in all TUs
 - **Ordered** - all other static storage duration variables in order of definition in TU
- Implementation defined if happens before `main()` or deferred until used (like static local variables)

`basic.start.dynamic`



This is a simplification. Hot Take: avoid dynamic initialisation.

STATIC LOCAL VARIABLES

- Standard states initialisation strongly happens before first access; but implementation defined if happens before `main()`
- What really happens:
 - *dynamic initialisation* will happen first time control passes through declaration
 - Requires compiler generated locks for thread safety
 - **Unless** constant/zero initialised, which happens before `main()`

`basic.start.dynamic`

STATIC STORAGE COMPILER IMPLEMENTATION

- How does a compiler/linker make this work?
- Will look at *gcc* and *Clang*, but essentially same in all compilers
- Goal is to understand why initialisation is specified this way

Sorry, some assembly code follows...

CONSTANT INITIALISATION

example.cpp

```
1 // both static duration
2 int globalCounter{10};      // external linkage
3 static int localCounter{11}; // internal linkage
```

compiler-explorer.com/z/4bjoTnGqY

Values are already set, after program is loaded before any code runs because it is written directly into the program file

x86-64 Asm - gcc

```
1      .globl globalCounter    ; export symbol
2      .data                  ; use .data segment
3 globalCounter:
4      .long 10                ; globalCounter=10
5
6 localCounter:
7      .long 11                ; localCounter=11
```


ZERO INITIALISATION

example.cpp

```
1 // both static duration
2 int defaultInitInt;
3 int zeroInitInt{0};
4
5 void *defaultInitPtr;
6 int *nullInitPtr{nullptr};
```

compiler-explorer.com/z/bn7T4Po99

Values are zeroed automatically by the loader before any code runs. The **.bss** segment is not actually stored in the program file, just the required size.

x86-64 Asm - gcc

```
1      .globl  nullInitPtr      ; export symbol
2      .bss                                ; use .bss segment
3  nullInitPtr:
4      .zero   8                  ; nullInitPtr = 0/nullptr
5
6      .globl  defaultInitPtr   ; export symbol
7  defaultInitPtr:
8      .zero   8                  ; defaultInitPtr = 0/nullptr
9
10     .globl  zeroInitInt; export symbol
11  zeroInitInt:
12     .zero   4                  ; zeroInitInt = 0
13
14     .globl  defaultInitInt   ; export symbol
15  defaultInitInt:
16     .zero   4                  ; defaultInitInt = 0
```

i re-ordered by alignment?

DYNAMIC INITIALISATION

example.cpp

```
1 extern int externalInteger;  
2 int globalInteger1{externalInteger};  
3 int globalInteger2{externalInteger};
```

compiler-explorer.com/z/ja8G3b3jd

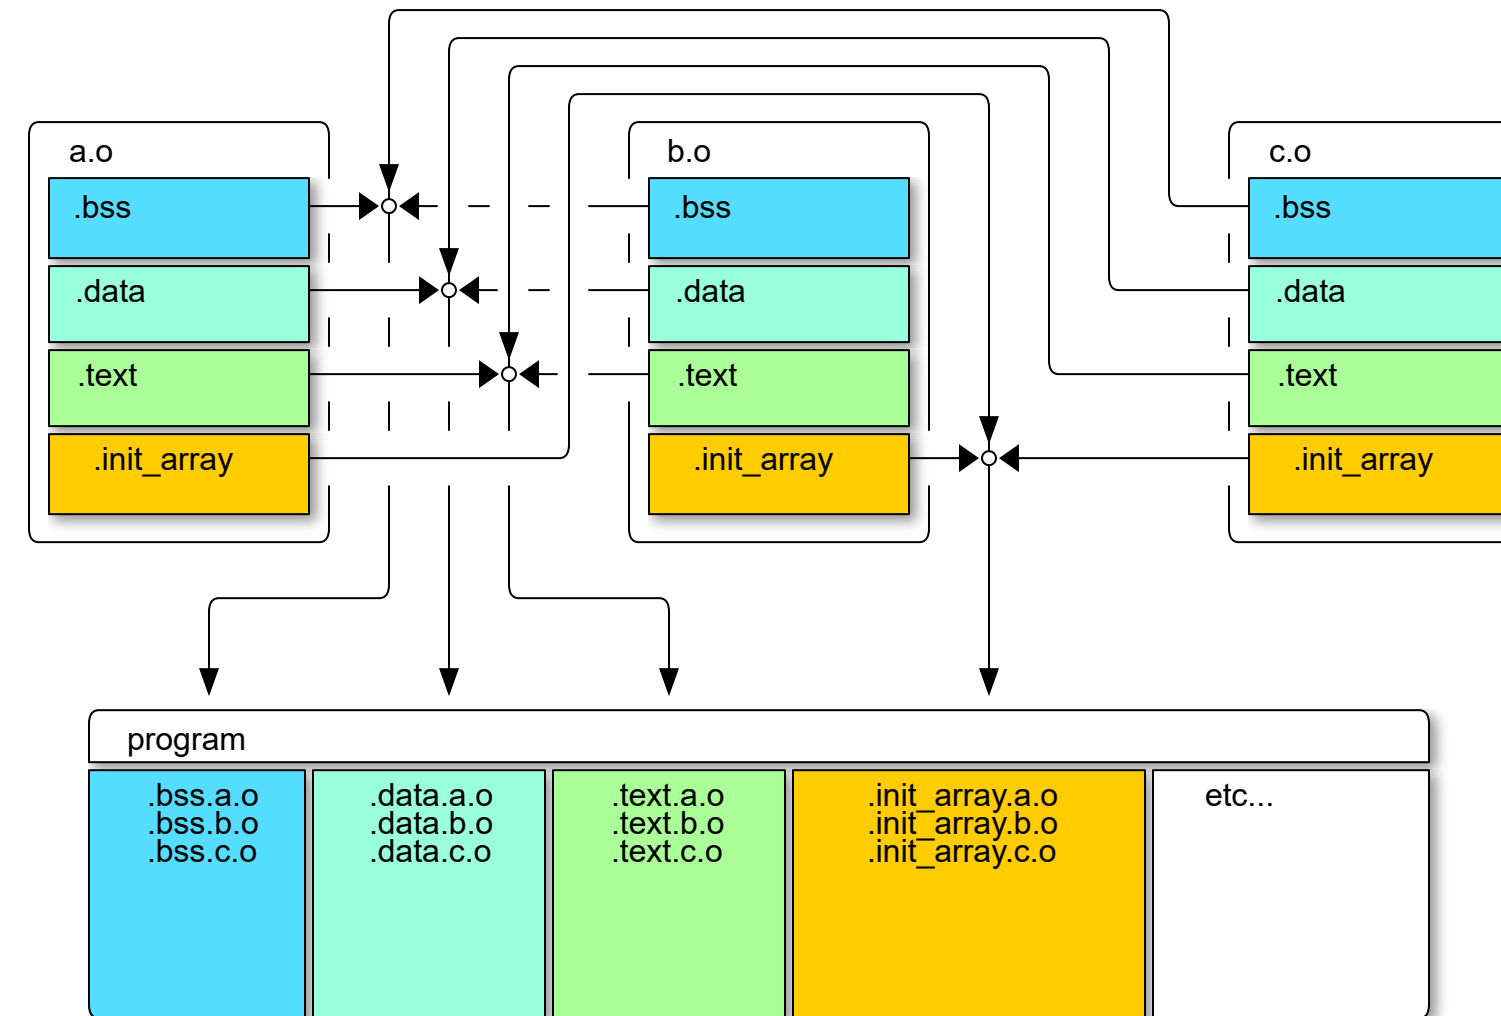
- Compiler can't determine the value
- Zero initialises the variables
- Writes function to compute value at run time. May register destructor with `std::atexit()`
- Adds that function to `.init_array` list to be called by startup code

x86-64 Asm - Clang

```
1      .text                ; use code segment  
2  _GLOBAL__sub_I_example.cpp:  
3      mov     rax, qword ptr [rip + externalInteger@GOTPCREL]  
4      mov     eax, dword ptr [rax]  
5      ; Initialise variables, in order from TU  
6      mov     dword ptr [rip + globalInteger1], eax  
7      mov     dword ptr [rip + globalInteger2], eax  
8      ret  
9  
10     .bss                 ; use .bss segment  
11     .globl  globalInteger1 ; export symbol  
12  globalInteger1:  
13     .long   0  
14  
15     .globl  globalInteger2 ; export symbol  
16  globalInteger2:  
17     .long   0  
18  
19     ; place the address of _GLOBAL__sub_I_example into  
20     ; a section called .init_array  
21     ; Startup code calls all these before main()  
22     .section .init_array,"aw",@init_array  
23     .quad   _GLOBAL__sub_I_example.cpp
```

DETOUR: LINKING

- Linker gathers all the sections and puts them together in program



- Order of object file sections depends on link order
- Link order depends on ??? (build system, configuration, etc..)

DETOUR: START UP `.init_array`

Artist's impression

- Implementation supplied start up code:
 - Establishes environment
 - Runs initialisers in `.init_array`
 - calls `main()`
- Remember, initialiser order is fragile/unpredictable!
 - Dynamic initialisation can happen in unpredictable order!
- `std::atexit()` has list of calls to run after `main()`
 - Reverse order of registration

```
1 // pointer to function returning void and taking
2 // no arguments.
3 using void_fn = void(*)(void);
4
5 // provided by linker from segment bounds
6 extern void_fn *init_array_start;
7 extern void_fn *init_array_end;
8
9 extern void setup_environment();
10 extern int main();
11
12 std::vector<void_fn> exit_list{};
13 void atexit(void_fn f) { exit_list.push_back(f); }
14
15 int start()
16 {
17     setup_environment();
18
19     std::span init{init_array_start, init_array_end};
20     for(auto *fn : init) {
21         fn();
22     }
23
24     auto exit = main();
25
26     for(auto *fn : exit_list | std::ranges::views::reverse) {
27         fn();
28     }
29
30     return exit;
31 }
```

STATIC LOCAL VARIABLES - CONSTANT INITIALISATION

example.cpp

```
1 int& get_constant()
2 {
3     static int s_const{42}; // or 0
4     return s_const;
5 }
```

compiler-explorer.com/z/q94rh66Pd

- Same as constant initialisation of any other variable
- Also works for zero initialisation

x86-64 Asm - Clang

```
1     .text                ; use .text (code) segment
2     .globl get_constant() ; export function
3 get_constant():
4     lea     rax, [rip + get_constant()::s_const]
5     ret
6
7     .data                ; use .data segment
8 get_constant()::s_const:
9     .long   42            ; get_constant()::s_const = 42
```

```
1     ; get_constant() as above
2
3     .bss                ; use .bss segment
4 get_constant()::s_const:
5     .long   0            ; get_constant()::s_const = 0
```

STATIC LOCAL VARIABLES - DYNAMIC INITIALISATION

example.cpp

```
1 extern int externalValue;
2
3 int& get_dynamic()
4 {
5     static int s_dyn{42*externalValue};
6     return s_dyn;
7 }
```

compiler-explorer.com/z/5zMbac9cE

- Effectively initialises first time control passes the definition
- Compiler creates a guard variable
- Double check locking
- Costs single check even after initialised
- `__cxa_guard_acquire` and `__cxa_guard_release` generally involve global mutex operation

x86-64 Asm - Clang

```
1     .text
2     .globl get_dynamic() ; export function
3 get_dynamic():
4     movzx    eax, byte ptr [rip + guard for get_dynamic()::s_dyn]
5     test     al, al      ; check if already initialised
6     je       .LBB0_1
7     lea      rax, [rip + get_dynamic()::s_dyn]
8     ret      ; return the value
9
10    .LBB0_1:
11    push     rax
12    lea      rdi, [rip + guard for get_dynamic()::s_dyn]
13    call     __cxa_guard_acquire@PLT
14    test     eax, eax     ; double check if initialised
15    je       .LBB0_3     ; skip if it was
16    mov      rax, qword ptr [rip + externalValue@GOTPCREL]
17    imul     eax, dword ptr [rax], 42
18    mov      dword ptr [rip + get_dynamic()::s_dyn], eax
19    lea      rdi, [rip + guard for get_dynamic()::s_dyn]
20    call     __cxa_guard_release@PLT
21    .LBB0_3:
22    add      rsp, 8
23    lea      rax, [rip + get_dynamic()::s_dyn]
24    ret      ; return value
25
26    ; local .data segment allocations
27    .local  get_dynamic()::s_dyn
28    .comm   get_dynamic()::s_dyn,4,4
29    .local  guard for get_dynamic()::s_dyn
30    .comm   guard for get_dynamic()::s_dyn,8,8
```

CONSTRUCTORS AND DESTRUCTORS FOR STATIC STORAGE DURATION

- Dynamic Initialisation for Non Trivial types (with constructors and destructors)
- With possible exception for `constexpr` constructors
- The destructor will be registered with `std::atexit()` when constructed

CONSTRUCTOR AND DESTRUCTOR EXAMPLE

example.cpp

```
1 struct S
2 {
3     S();
4     ~S();
5 };
6
7 S s;
```

- Familiar Dynamic Initialisation pattern calls **ctor**
 - Registers **dtor** with **std::atexit()** (**__cxa_atexit**)
- Similar with local variables of static duration
 - Registers **dtor** after construction

compiler-explorer.com/z/1af88d3s6

x86-64 Asm - Clang

```
1      .text                                ; use code segment
2  _GLOBAL__sub_I_example.cpp:
3      push    rbx
4      ; call constructor
5      lea     rbx, [rip + s]
6      mov     rdi, rbx
7      call    S::S()@PLT
8      ; register destructor
9      mov     rdi, qword ptr [rip + S::~~S()@GOTPCREL]
10     lea     rdx, [rip + __dso_handle]
11     mov     rsi, rbx
12     pop     rbx
13     jmp     __cxa_atexit@PLT ; tail call
14
15     .bss                                ; use .bss segment
16     .globl  s                            ; export symbol
17 s:
18     .zero   1                            ; must have unique address
19
20     ; place the address of _GLOBAL__sub_I_example into
21     ; a section called .init_array
22     ; Startup code calls all these before main()
23     .section .init_array,"aw",@init_array
24     .quad   _GLOBAL__sub_I_example.cpp
```


SO, WHAT DOES THIS MEAN?

- Constant or Zero initialisation is predictable and reliable
- Dynamic initialisation is much harder:
 - Different types of initialisation, even in the same file, may not be predictably ordered
 - All bets are off across TUs
 - The linker will choose
 - The order will likely be based on argument order in the link command
 - The order can change arbitrarily
 - It *might* work now, but can easily change

This is where the *Static Initialisation Order Fiasco* comes from!

COMMON INITIALIZATION PATTERNS

- Meyers' Singleton (Construct on first use)
- Nifty Counter
- Static buffer and Reference

MEYERS' SINGLETON

- Very common Singleton pattern
- Can be free function or class member
- Initialisation on first use

```
1 class Singleton
2 {
3 public:
4     static Singleton& instance()
5     {
6         static Singleton instance{};
7         return instance;
8     }
9
10 private:
11     Singleton();
12     ~Singleton();
13 };
14
15 class log { /* ... */ };
16
17 log& get_log()
18 {
19     static log theLog{};
20     return theLog;
21 }
```

DEPENDENT SINGLETONS?

Looks Good; ship it!

```
1 struct Log
2 {
3     static Log& get() {
4         static Log instance{}; return instance;
5     }
6
7     void write(std::string_view msg) { fmt::print("LOG: {}\n", msg); }
8
9     Log() { fmt::print("Log::Log()\n"); }
10    ~Log() { fmt::print("Log::~~Log()\n"); }
11 };
12
13 struct Singleton
14 {
15     static Singleton& get() {
16         static Singleton instance{}; return instance;
17     }
18
19     void test() { Log::get().write("Singleton::test()"); }
20
21     Singleton() {
22         fmt::print("Singleton::Singleton()\n");
23         Log::get().write("in Singleton ctor");
24     };
25     ~Singleton() {
26         fmt::print("Singleton::~~Singleton()\n");
27         Log::get().write("in Singleton dtor");
28     };
29 };
30
31 int main() { Singleton::get().test(); }
```

```
1 Singleton::Singleton()
2 Log::Log();
3 LOG: in Singleton ctor
4 LOG: Singleton::test();
5 Singleton::~~Singleton()
6 LOG: in Singleton dtor
7 Log::~~Log();
```

compiler-
[explorer.com/z/cv53nq7cz](https://godbolt.org/z/cv53nq7cz)

DEPENDENT SINGLETONS...

```
1 struct Log
2 {
3     static Log& get() {
4         static Log instance{}; return instance;
5     }
6
7     void write(std::string_view msg) { fmt::print("LOG: {}\n", msg); }
8
9     Log() { fmt::print("Log::Log()\n"); }
10    ~Log() { fmt::print("Log::~~Log()\n"); }
11 };
12
13 struct Singleton
14 {
15     static Singleton& get() {
16         static Singleton instance{}; return instance;
17     }
18
19     void test() { Log::get().write("Singleton::test()"); }
20
21     Singleton() {
22         fmt::print("Singleton::Singleton()\n");
23         /* *** Log::get().write("in Singleton ctor"); *** */
24     };
25     ~Singleton() {
26         fmt::print("Singleton::~~Singleton()\n");
27         Log::get().write("in Singleton dtor");
28     };
29 };
30
31 int main() { Singleton::get().test(); }
```

```
1 Singleton::Singleton()
2 Log::Log();
3 LOG: Singleton::test();
4 Log::~~Log();
5 Singleton::~~Singleton()
6 LOG: in Singleton dtor
```

- **Log** use after destruction!
- **std::atexit()** ordering problem
- Worked because **Log** was constructed before **Singleton** finished construction
- Now, **Log** is constructed after **Singleton**



Avoid Dependencies -
Especially in ctor/dtor

NIFTY COUNTER

- Allows us to take direct control over lifetime
- Injects an *Internal Linkage* object of *Static Storage* duration into each TU that includes the header
- Injection happens before any use of dependencies - Ordered init in TU
- Nifty counter object `ctor/dtor` maintain reference count
- first `ctor` call constructs dependencies
- last `dtor` call destroys dependencies
- `<iostream>` uses this for `std::cout`, `std::cin` etc.

NIFTY COUNTER EXAMPLE

nifty.hpp

```
1 #include <vector>
2
3 static struct nifty_counter {
4     nifty_counter();
5     ~nifty_counter();
6 } const NiftyCounterInstance;
7
8 // the global resource
9 extern std::vector<int> *someVector;
```

main.cpp

```
1 #include "nifty.hpp"
2 #include <vector>
3
4 int main() {
5     return someVector->size();
6 }
```

nifty.cpp

```
1 #include "nifty.hpp"
2 #include <vector>
3
4 static int instance_count{0};           // zero init
5 std::vector<int> *someVector{nullptr}; // zero init
6
7 nifty_counter::nifty_counter() {
8     if(instance_count++ == 0) {
9         someVector = new std::vector{1, 2, 3, 4};
10    }
11 }
12
13 nifty_counter::~nifty_counter() {
14     if(--instance_count == 0) {
15         delete someVector;
16         someVector = nullptr;
17     }
18 }
```

compiler-explorer.com/z/v4K5rqjhM

STATIC BUFFER & REFERENCE

- Allows a global reference to object
- Reserve a buffer to create object within
- Reference is constant initialised to point to buffer
- use *Nifty Counter* to control initialisation

EXAMPLE

log.hpp

```
1 struct Log {
2     Log();
3     ~Log();
4     int test() { return 42;};
5 };
6 extern Log& log;
7
8 static struct LogInitialiser {
9     LogInitialiser();
10    ~LogInitialiser();
11 } const logInitialiser;
```

main.cpp

```
1 #include "log.hpp"
2
3 int main() {
4     return log.test();
5 }
```

compiler-explorer.com/z/75Yn5jTqr

log.cpp

```
1 #include <new>           // placement new
2 #include <type_traits>    // aligned_storage
3
4 // internal linkage
5 static int init_counter;
6 static std::aligned_storage_t<sizeof(Log), alignof(Log)> logBuf;
7
8 // external linkage
9 Log& log = reinterpret_cast<Log&>(logBuf);
10
11 Log::Log() {}
12 Log::~~Log() {}
13
14 LogInitialiser::LogInitialiser() {
15     if (init_counter++ == 0) {
16         new (&log) Log();
17     }
18 }
19
20 LogInitialiser::~~LogInitialiser() {
21     if (--init_counter == 0) {
22         log.~Log();
23     }
24 }
```

LINKING AND LOADING LIBRARIES

- What does the C++ standard say about linking and loading libraries?
 - Nothing - just the *Abstract Machine*
 - Everything else is *Implementation Defined*

STATICALLY LINKED LIBRARIES

- Just a container of segments and symbols `.a` / `.lib`
- Linker adds them to the executable
- This includes the `.data` and `.bss` segments for constant/zero initialisation
- Also the `.init_array` segment for dynamic initialisation
- Therefore same ordering issues

SHARED LIBRARIES (DYNAMICALLY LINKED LIBRARIES)

- Linker flags that the `.so` / `.dll` should be loaded at run time
- Program loader does this and wires up symbols
- Compiler must ensure that initialisation occurs **as if** it was running on the *Abstract Machine*
- Program Loader doesn't know anything about C++ rules
- Can be loaded/unloaded at run time as well

PLATFORM SPECIFIC HOOKS

⚠ If you are using this, you're well outside of C++ standard — beyond the scope of this session!

- Linux
 - `__attribute__((constructor))` / `__attribute__((destructor))`
 - Any number of functions placed in `.ctors` & `.dtors` segments — just like `.init_array`
 - Ordering of calls is undefined
 - `_init()` and `_fini()` can be used to run specific code
- Windows
 - Must run all initialisers on load, finalisers on unload
 - `DllMain()` is called on process / thread attach / detach to allow you to run specific code

LIBRARY STATIC INITIALISATION

- How can we ensure proper a safe initialisation in a library?
 - What can we rely on?
 - Cross platform?

WHAT CAN WE RELY ON?

- Zero Initialisation
- Constant Initialisation
- ~~Dynamic Initialisation~~
- Ordered initialisation within each TU
- Nifty Counter
- Meyers' Singleton - without interdependencies
- Class Member initialisation order

SOME IDEAS

- Several approaches are possible
- A simple library might get away with a simple approach
- Combinations for more complex scenarios

IDEA 1: CLIENT CALLS FUNCTIONS

Provide `start()` and `stop()` methods that client calls to initialise state and clean up

PROS:

- Simple & Traditional - many existing C libraries
- Allocate and initialise (library) global state in `start()`
- Destroy and de-allocate (library) global state in `stop()`
- Can inject configuration into start up

CONS:

- Easy for client to forget to call one or both
- Generally run-time errors
- Is it valid to `start()` and `stop()` multiple times?
- What about threads?

IDEA 2: NIFTY COUNTER TO INITIALISE

Use a Nifty Counter to automatically call `start()` and `stop()` like functions

PROS:

- Automatic by including library header
- STL uses this for `<iostream>`
- Initialisation happens before `main()`

CONS:

- Can only initialise before `main()` and destroy after - program lifetime
- No opportunity for configuration to be injected

IDEA 3: ONE SINGLETON TO RULE THEM ALL

Wrap all (library) global state into a container object and Access it through Meyers' Singleton

PROS:

- Nothing for client to do
- All initialisation happens in one place - avoid singleton interdependencies
- Container class can order initialisation and pass dependencies between members

CONS:

- Lacks precise control on timing
- No opportunity for configuration to be injected
- Can couple many parts of the library to one TU (can be mitigated)

IDEA 4: RAI LIBRARY HANDLE

- Library exposes top level handle; client controls lifetime explicitly
- All Api calls go through the handle object
- Manages (library) global state lifetime

PROS:

- Client can choose to have one or many non-overlapping lifetimes
- Intuitive if all (top level) calls are handle member functions
- Can inject configuration into start up
- Can also create non-static dependencies, threads, queues etc.

CONS:

- Client can store returned objects outside of lifetime scope
 - Use-after-free: prefer Value Semantics
 - Not specific to this option
- Can couple many parts of the library to one TU (use compilation firewalls — PIMPL)

LIBRARY EXAMPLE IMPLEMENTATION

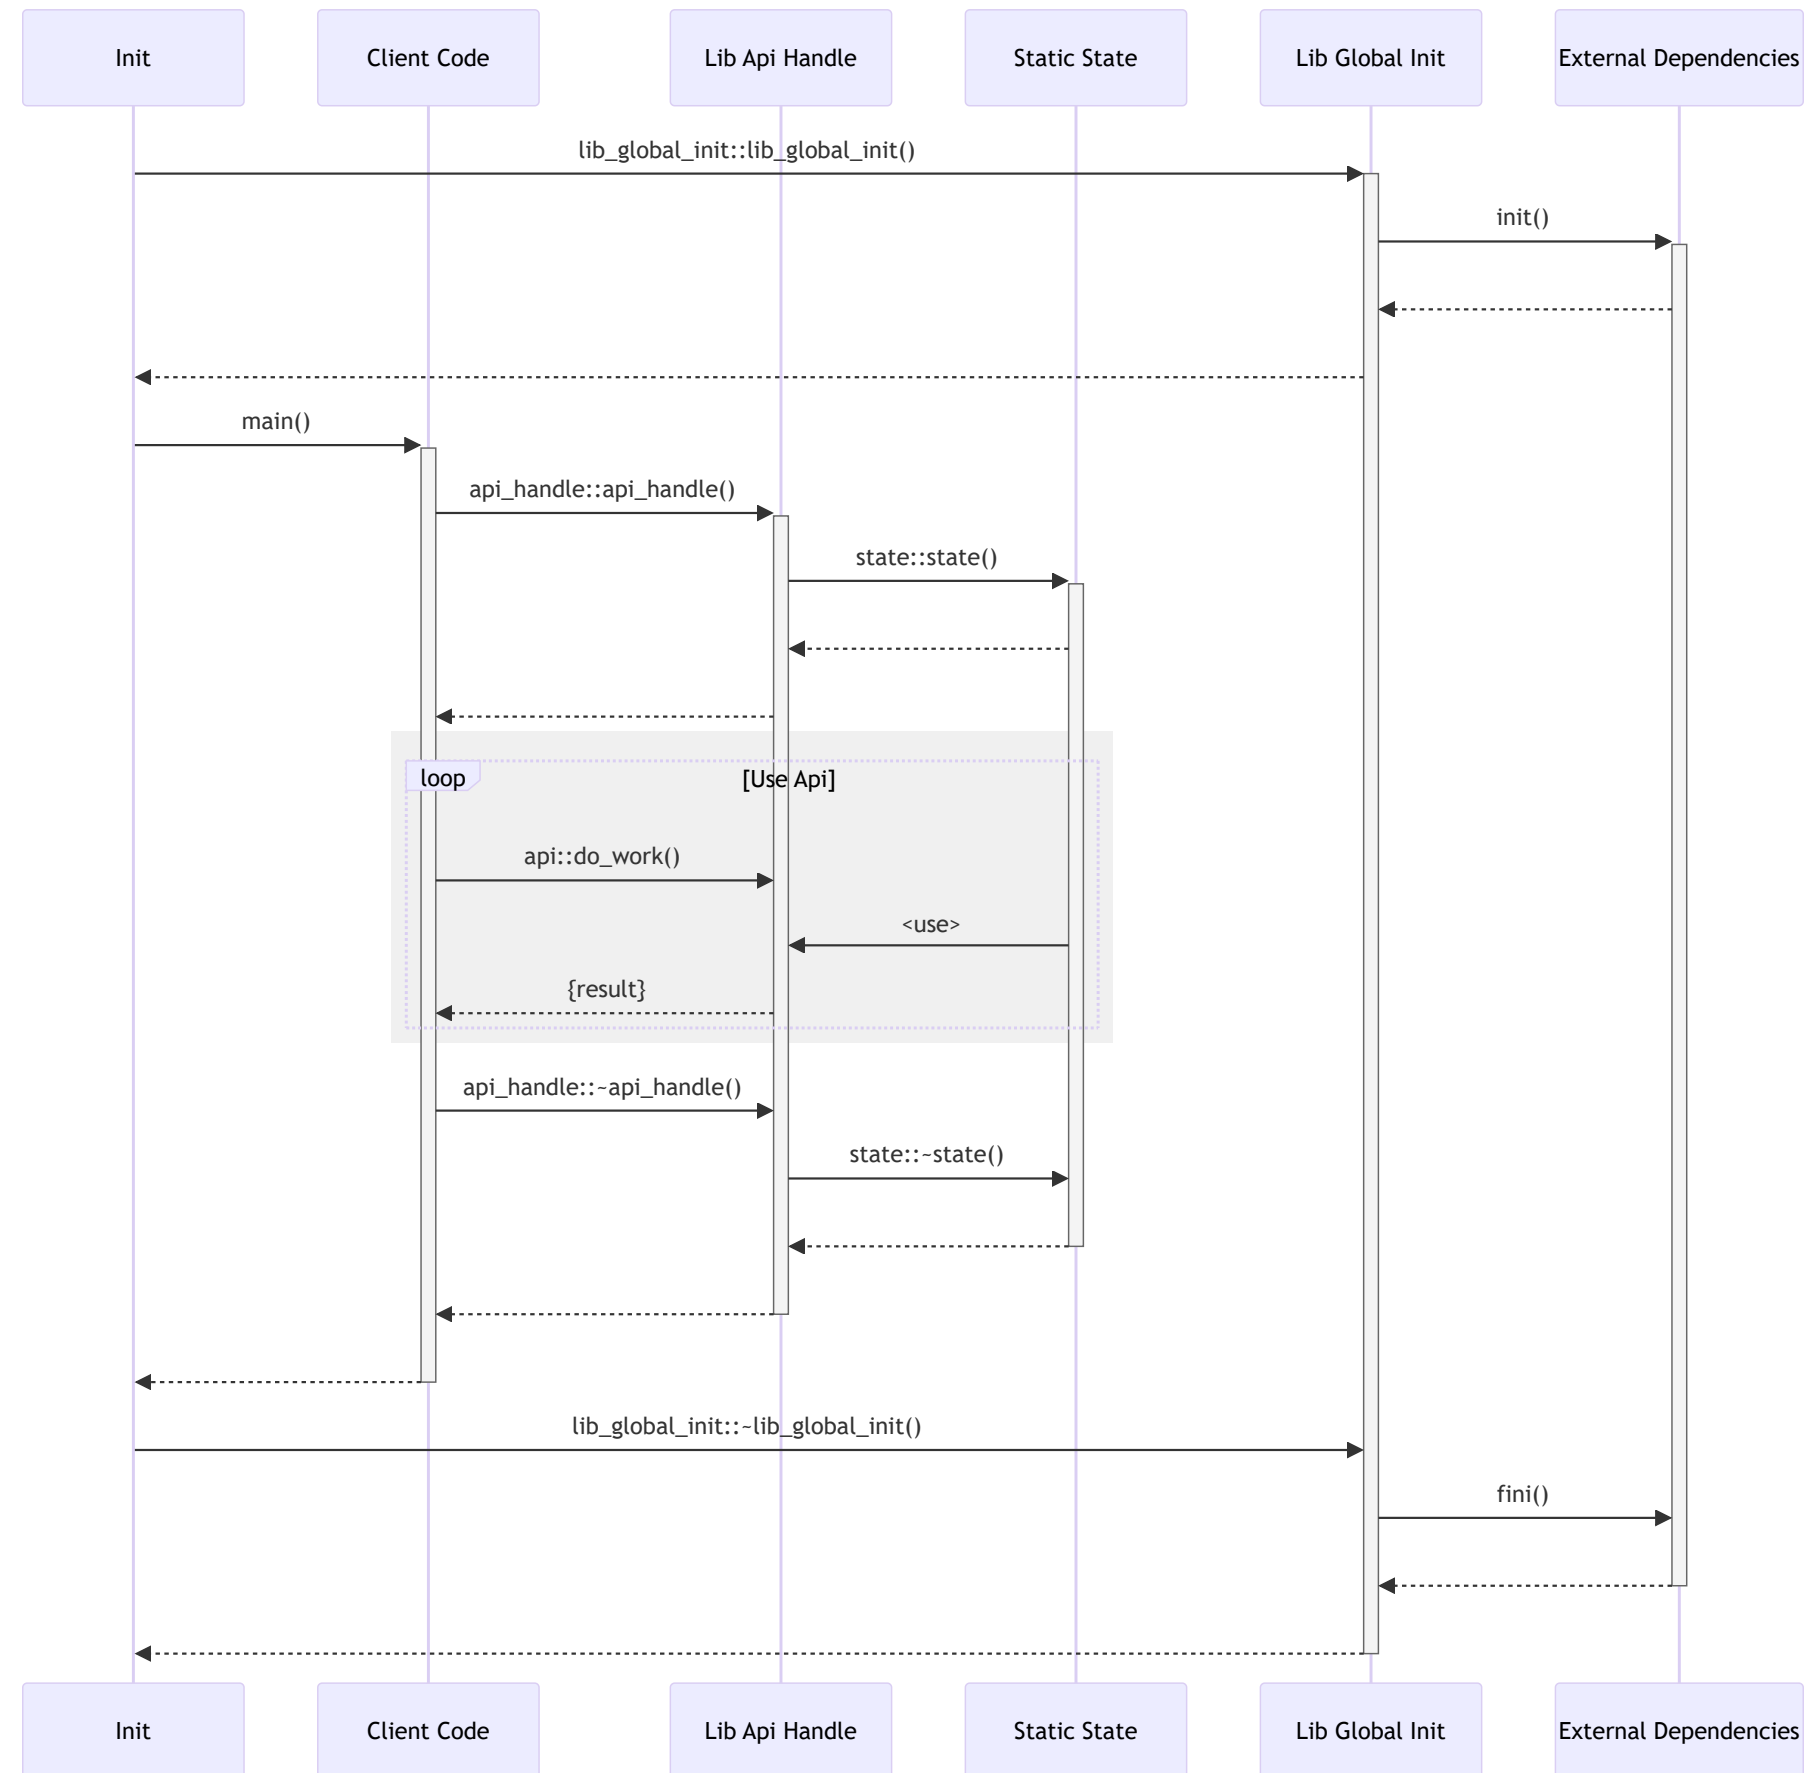
Walk-through of a worked example, imaginatively named `libExample`.

- Uses RAII Handle Object to manage lifetimes with single instance validation
- Uses Nifty Counter for one-time initialisation
- Uses PIMPL idiom for compilation fire-walling implementation
- Separates lifetime management from (internal) global access

Code is simplified for slides, but full working example:

 github.com/AshleyRoll/conferences/cppnorth2024/code

LIBRARY EXAMPLE LIFETIME



LIBEXAMPLE FUNCTIONALITY

- Lets have a worker thread that pulls items from a queue
- The client adds work to the queue
- There is a logger that is accessible to the worker and other components
- Not shown in presentation, but available in code example
 - `queue` is a simple thread safe queue which provides a producer and consumer object
 - `worker` is a wrapper over `std::jthread` which takes a queue consumer

LIBEXAMPLE API - RAII OBJECT

example_api.hpp

```
1 // forward declare PIMPL abstraction for example_api
2 class example_api_impl;
3
4 // The public interface and RAII handle to the
5 // example library.
6 class example_api
7 {
8 public:
9     example_api();
10    ~example_api();
11
12    // disable copy and only allow move for the lifetime
13    // management of the API resources.
14    example_api(example_api &) = delete;
15    auto operator=(example_api const &)
16        -> example_api & = delete;
17    example_api(example_api &&) = default;
18    auto operator=(example_api &&)
19        -> example_api & = default;
20
21    // Expose public API here, and delegate to the
22    // PIMPL class internally
23    void do_work(std::string const &message);
24
25 private:
26    std::unique_ptr<example_api_impl> m_impl; // PIMPL
27 };
```

example_api.cpp

```
1 example_api::example_api()
2     : m_impl{ std::make_unique<example_api_impl>() }
3 {
4     fmt::println("example_api::ctor()");
5 }
6
7 example_api::~example_api()
8 {
9     fmt::println("example_api::~dtor()");
10 }
11
12 void example_api::do_work(std::string const &message)
13 {
14     m_impl->do_work(message);
15 }
```


LIBEXAMPLE API IMPLEMENTATION

example_api_impl.hpp

```
1 class example_api_impl
2 {
3 public:
4     example_api_impl();
5     ~example_api_impl();
6
7     // TODO: Implement public API here
8     void do_work(std::string const &message);
9
10    // disable move and copy
11    example_api_impl(example_api_impl &) = delete;
12    auto operator=(example_api_impl const &)
13        -> example_api_impl & = delete;
14    example_api_impl(example_api_impl &&) = delete;
15    auto operator=(example_api_impl &&)
16        -> example_api_impl & = delete;
17
18 private:
19     library_instance_lock m_instanceLock{};
20     scoped_singleton<logger> m_log;
21     queue m_workQueue{};
22     worker m_worker;
23 };
```

example_api_impl.cpp

```
1 example_api_impl::example_api_impl()
2     : m_log{ "static_log" }
3     , m_worker{ m_workQueue.get_consumer() }
4 {
5     fmt::println("example_api_impl::ctor()");
6 }
7
8 example_api_impl::~example_api_impl()
9 {
10    fmt::println("example_api_impl::~dtor()");
11 }
12
13 void example_api_impl::do_work(
14     std::string const &message)
15 {
16     // log shortcut, we will see how shortly
17     log::info(
18         fmt::format(R"(example_api_impl::do_work("{}"))",
19             message));
20     m_workQueue.get_producer().send(message);
21 }
```

SINGLE ACTIVE HANDLE

- Validate only a single active handle before any other initialisation
- If this fails, the `example_api_impl` ctor will fail and no other members (`m_log`, `m_workQueue`, `m_worker`) will be created.
- `constexpr` (C++20) ensures our static storage is constant/zero initialised
 - compiler will error if it has to resort to dynamic initialisation

`library_instance_lock.hpp`

```
1 struct library_instance_lock
2 {
3     constexpr inline static std::atomic_bool s_instanceAlive{};
4
5     library_instance_lock()
6     {
7         bool expected{ false };
8         if (not s_instanceAlive.compare_exchange_strong(expected, true)) {
9             throw library_instance_already_exists();
10        }
11    }
12
13    ~library_instance_lock()
14    {
15        s_instanceAlive.store(false);
16    }
17
18    // ... disable move/copy ...
19 };
```

SEPARATE LIFETIME FROM ACCESS

scoped_singleton.hpp

- `scoped_singleton<T>`
- Separates the control of lifetime from access
 - Access is effectively a global pointer
 - Lifetime is managed by owning object - `example_api_impl`

```
1 template<typename T>
2 class scoped_singleton
3 {
4 public:
5     template<typename... Args>
6     explicit scoped_singleton(Args &&...args)
7         : m_instance(std::forward<Args>(args)...)
8     {
9         if (s_instance != nullptr) {
10             throw std::logic_error(
11                 "scoped_singleton<T> is already active");
12         }
13         s_instance = &m_instance;
14     }
15
16     ~scoped_singleton() { s_instance = nullptr; }
17
18     // ... disable move/copy - static pointer would be broken
19
20     static auto instance() -> T &
21     {
22         if (s_instance == nullptr) {
23             throw std::logic_error("no live instance");
24         }
25         return *s_instance;
26     }
27
28 private:
29     T m_instance;
30     constexpr inline static T *s_instance{nullptr};
31 };
```

ACCESSING SCOPED_SINGLETON<LOGGER>

- The `scoped_singleton` provides a reference to the active `instance()`
- A bit much to type everywhere, so wrap it in some helper functions

```
1 class logger
2 {
3 public:
4     explicit logger(std::string name);
5     ~logger();
6
7     void error(std::string const &message);
8     void warn(std::string const &message);
9     void info(std::string const &message);
10
11 private:
12     std::string m_name;
13 };
```

```
1 namespace log {
2     inline static void error(std::string const &message)
3     {
4         scoped_singleton<logger>::instance().error(message);
5     }
6
7     inline static void warn(std::string const &message)
8     {
9         scoped_singleton<logger>::instance().warn(message);
10    }
11
12    inline static void info(std::string const &message)
13    {
14        scoped_singleton<logger>::instance().info(message);
15    }
16 } // namespace log
```

ONE TIME INITIALIZATION

- What if `libExample` depended on 3rd party libraries?
 - `libcurl` or `openssl` require one time initialisation and cleanup
- This is independent on the lifetime of our `example_api` RAI object
- Use a *Nifty Counter* to initialise automatically

example_api.hpp

```
1 static struct library_global_initializer
2 {
3     library_global_initializer() noexcept;
4     ~library_global_initializer() noexcept;
5
6     // ... disable move/copy
7     // - avoid corrupting counter ...
8 } const LibraryGlobalInitializer;
```

library_global_initializer.cpp

```
1 namespace {
2     void global_initialize() noexcept
3     {
4         // TODO: Call all dependencies start up routines here
5         // Example:
6         //     libcurl - curl_global_init()
7         fmt::println("Initializing lib::example dependencies");
8     }
9
10    void global_finalize() noexcept
11    {
12        // TODO: Call all dependencies clean up routines here
13        // Example:
14        //     libcurl - curl_global_cleanup()
15        fmt::println("Finalizing lib::example dependencies");
16    }
17
18    int initializationCounter;
19 } // namespace
20
21 library_global_initializer::library_global_initializer() noexcept
22 {
23     // if this is the first call, we run the initialization
24     if (initializationCounter++ == 0) { global_initialize(); }
25 }
26
27 library_global_initializer::~library_global_initializer() noexcept
28 {
29     // if this is the last call, we run the finalization
30     if (--initializationCounter == 0) { global_finalize(); }
31 }
```

CLIENT CODE

client.cpp

```
1 int main() {
2   example_api api{}; // api lifetime begins
3
4   api.do_work("task 1");
5   api.do_work("task 2");
6
7 } // api lifetime ends
```

```
1 int main() {
2   example_api api{}; // api lifetime begins
3
4   api.do_work("task 1");
5   api.do_work("task 2");
6
7   // ...
8
9   example_api another{}; // Exception
10  another.do_work("nope");
11
12 } // api lifetime ends
```

ADDITIONAL NICETIES

The example code project also:

- Verifies that the client code links against the same version of the code the header is from
- CMake project with presets for
 - `clang` and `gcc`
 - *static* vs *shared* library
 - Ensures symbols are hidden by default in the library

THANK YOU, QUESTIONS?

GitHub  <https://github.com/AshleyRoll>

Discord  *#include<C++>* @AshleyRoll

Slack  *Cpplang* @AshleyRoll